

Tutorial for Diffusion Model

What is Diffusion Model (DDPM)?

Denoising Diffusion Probabilistic Models

Jonathan Ho
UC Berkeley
jonathanho@berkeley.edu

Ajay Jain
UC Berkeley
ajayj@berkeley.edu

Pieter Abbeel
UC Berkeley
pabbeel@cs.berkeley.edu

[Denoising diffusion probabilistic models](#)

[\[PDF\] neurips.cc](#)

[J Ho](#), [A Jain](#), [P Abbeel](#) - Advances in neural information ..., 2020 - proceedings.neurips.cc

... This paper presents progress in [diffusion probabilistic models](#) [53]. A [diffusion probabilistic model](#) (which we will call a “[diffusion model](#)” for brevity) is a parameterized Markov chain ...

☆ 保存 99 引用 被引用次数: 28666 相关文章 [所有 10 个版本](#) 99

What is Diffusion Model?

- The diffusion model is a class of generative models widely used for images, videos, and audio.




Figure 1: Generated samples on CelebA-HQ 256×256 (left) and unconditional CIFAR10 (right)

How does Diffusion Model work?

- Diffusion Model works by gradually adding random noise to data through a forward diffusion process, and then learning to reverse this process to reconstruct or generate realistic samples from pure noise.



Figure 6: Unconditional CIFAR10 progressive generation ($\hat{x}_0$ over time, from left to right). Extended samples and sample quality metrics over time in the appendix (Figs. 10 and 14).

How does Diffusion Model work?

- Goal: Suppose we have 100 different pictures of dogs (denoted as X_0), and we want to generate a new picture of a dog that looks realistic but different from any of the existing ones.
- Step 1: Add a little bit of noise to the original picture of dogs (denoted as X_1)
- Step 2: Add a little bit of noise to X_1 (denoted as X_2)
- Step 3: Repeat this process until we have a picture of noise
- Step 4: Train a neural network to denoise from X_t to X_{t-1} until X_0
- Step 5: Now we have a neural network that can start from pure noise and denoise it step by step to generate a new, realistic picture of a dog.



A taste of Math

- Stochastic Process
 - We add random noise to images. So the sequence $(x_0, x_1, \dots, x_t)$ is a stochastic process, and we use $q(x_t)$ to describe the probabilistic distribution of x_t within this process.
 - In a Stochastic process, we often optimize expectations with respect to probability distributions.
- Gaussian Distribution (Also known as Normal Distribution)
 - A Gaussian (or Normal) distribution often arises as the sum of many independent random variables.
 - Since real-world noise can usually be modeled as the combination of many small random effects, it is common to assume that the added noise follows a Gaussian distribution. Thus, we use Gaussian noise in diffusion models.

A taste of Math

- Conditional Probability

- $q(x_t)$ represents the (marginal) probability distribution of x_t .
- When the original clean image x_0 is known, $q(x_t|x_0)$ denotes the conditional probability distribution of x_t given x_0

- KL Divergence

- The Kullback–Leibler (KL) divergence measures how one probability distribution differs from another.

$$D_{\text{KL}}(P \parallel Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$


Key Insights Behind Diffusion Models

1. Markov Property

- each noisy image $\mathbf{x}_t$ depends only on the previous step $\mathbf{x}_{t-1}$, not on earlier steps.
- $q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_{t-2}, \dots, \mathbf{x}_0) = q(\mathbf{x}_t | \mathbf{x}_{t-1})$


2. Forward Diffusion Process (mathematical definition of “adding noise”)

- $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$
- $q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$
- $q(\mathbf{x}_{1:T} | \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})$
- $q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$ using the notation $\alpha_t := 1 - \beta_t$ and $\bar{\alpha}_t := \prod_{s=1}^t \alpha_s$



Key Insights Behind Diffusion Models

- 3. Reverse Diffusion Process (“denoising”) and Variational Inference
 - We aim to train a neural network (U-net) to perform denoising, in other words, to approximate $q(\mathbf{x}_{t-1} | \mathbf{x}_t)$.
 - During training, the original $\mathbf{x}_0$ is known, so the true posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$ can be computed analytically using the Bayesian rule. $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) \frac{q(\mathbf{x}_{t-1} | \mathbf{x}_0)}{q(\mathbf{x}_t | \mathbf{x}_0)}$
 - However, during generation, $\mathbf{x}_0$ is unknown. We employ variational inference by introducing a parameterized model $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ to approximate the true posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t)$, then minimizing the KL divergence.
 - Although $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$ is known during training, we still use $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ to approximate it, so that the learned model can later be used for denoising (generation) when $\mathbf{x}_0$ is not available.



Key Insights Behind Diffusion Models

$$\begin{aligned}
 L_{CE} &= -\mathbb{E}_{q(\mathbf{x}_0)} \log p_{\theta}(\mathbf{x}_0) \\
 &= -\mathbb{E}_{q(\mathbf{x}_0)} \log \left(\int p_{\theta}(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} \right) \\
 &= -\mathbb{E}_{q(\mathbf{x}_0)} \log \left(\int q(\mathbf{x}_{1:T}|\mathbf{x}_0) \frac{p_{\theta}(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \right) \\
 &= -\mathbb{E}_{q(\mathbf{x}_0)} \log \left(\mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \frac{p_{\theta}(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right) \\
 &\leq -\mathbb{E}_{q(\mathbf{x}_{0:T})} \log \frac{p_{\theta}(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \\
 &= \mathbb{E}_{q(\mathbf{x}_{0:T})} \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})} \right] = L_{VLB}
 \end{aligned}$$

• 4. Loss Function

- The loss function can be derived from the Evidence Lower Bound (ELBO) in Variational Inference.
- The exact mathematical form is theoretical and complicated. But we can understand it intuitively and rewrite it as three intuitive objectives.
- Diffusion loss just wants three things:

$$\begin{aligned}
 -\log p_{\theta}(\mathbf{x}_0) &\leq -\log p_{\theta}(\mathbf{x}_0) + D_{KL}(q(\mathbf{x}_{1:T}|\mathbf{x}_0) \| p_{\theta}(\mathbf{x}_{1:T}|\mathbf{x}_0)) && \text{; KL is non-negative} \\
 &= -\log p_{\theta}(\mathbf{x}_0) + \mathbb{E}_{\mathbf{x}_{1:T} \sim q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})/p_{\theta}(\mathbf{x}_0)} \right] \\
 &= -\log p_{\theta}(\mathbf{x}_0) + \mathbb{E}_q \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})} + \log p_{\theta}(\mathbf{x}_0) \right] \\
 &= \mathbb{E}_q \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})} \right] \\
 \text{Let } L_{VLB} &= \mathbb{E}_{q(\mathbf{x}_{0:T})} \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})} \right] \geq -\mathbb{E}_{q(\mathbf{x}_0)} \log p_{\theta}(\mathbf{x}_0)
 \end{aligned}$$

- At the beginning, The model's prior $p(\mathbf{x}_T)$ should match the true noisy distribution $q(\mathbf{x}_T|\mathbf{x}_0)$. $L_T = D_{KL}(q(\mathbf{x}_T|\mathbf{x}_0) \| p(\mathbf{x}_T))$
- In the middle, The model's reverse process $p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)$ should behave like the true diffusion process $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$.

$$L_t = D_{KL}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))$$

- At the end, the generated sample $\mathbf{x}_0$ should have no remaining noise.

$$L_0 = \mathbb{E}_{\mathbf{x}_0, t, \epsilon \sim \mathcal{N}(0, I)} [\|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|^2]$$

- Therefore, $L_{VLB} = L_T + L_{T-1} + \dots + L_0$

$$\begin{aligned}
 L_{VLB} &= \mathbb{E}_{q(\mathbf{x}_{0:T})} \left[\log \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{0:T})} \right] \\
 &= \mathbb{E}_q \left[\log \frac{\prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})}{p_{\theta}(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} \right] \\
 &= \mathbb{E}_q \left[-\log p_{\theta}(\mathbf{x}_T) + \sum_{t=1}^T \log \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1})}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} \right] \\
 &= \mathbb{E}_q \left[-\log p_{\theta}(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1})}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)} \right] \\
 &= \mathbb{E}_q \left[-\log p_{\theta}(\mathbf{x}_T) + \sum_{t=2}^T \log \left(\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} \cdot \frac{q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)} \right) + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)} \right] \\
 &= \mathbb{E}_q \left[-\log p_{\theta}(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} + \sum_{t=2}^T \log \frac{q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)} + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)} \right] \\
 &= \mathbb{E}_q \left[-\log p_{\theta}(\mathbf{x}_T) + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} + \log \frac{q(\mathbf{x}_T|\mathbf{x}_0)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \frac{q(\mathbf{x}_1|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)} \right] \\
 &= \mathbb{E}_q \left[\log \frac{q(\mathbf{x}_T|\mathbf{x}_0)}{p_{\theta}(\mathbf{x}_T)} + \sum_{t=2}^T \log \frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)}{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)} - \log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1) \right] \\
 &= \underbrace{\mathbb{E}_q [D_{KL}(q(\mathbf{x}_T|\mathbf{x}_0) \| p_{\theta}(\mathbf{x}_T))]}_{L_T} + \sum_{t=2}^T \underbrace{D_{KL}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} - \underbrace{\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}_{L_0}
 \end{aligned}$$

Key Insights Behind Diffusion Models

- 5. Simplified Loss Function

- Empirically, we can use a simplified loss function to replace the ELBO loss.
- Our neural network is trained to predict (remove) the added Gaussian noise at each step. In other words, we only need to let the predicted $\varepsilon_{\theta}(x_t, t)$ approximate the real noise ε .
- Therefore, the simplified and real-practical loss function is as follows.

$$L_{\text{simple}}(\theta) = \mathbb{E}_{t, x_0, \varepsilon} [\|\varepsilon - \varepsilon_{\theta}(x_t, t)\|^2]$$

Experiments



Figure 11: CelebA-HQ 256×256 generated samples




Figure 19: LSUN Cat generated samples. FID=19.75

Experiments




Figure 5: Unconditional CIFAR10 test set rate-distortion vs. time. Distortion is measured in root mean squared error on a $[0, 255]$ scale. See Table [4](#) for details.

Experiments



Figure 6: Unconditional CIFAR10 progressive generation ($\hat{x}_0$ over time, from left to right). Extended samples and sample quality metrics over time in the appendix (Figs. [10](#) and [14](#)).

Why do diffusion models work?

- Define the gradient of the data distribution (score function) as follows

$$s(x) = \nabla_x \log p(x),$$

- With the score, we can iteratively update a noisy sample toward the data using $x_{k+1} = x_k + \frac{\eta}{2} s(x_k) + \sqrt{\eta} z_k, \quad z_k \sim \mathcal{N}(0, I)$ (similar to Taylor expansion)
- Thus, diffusion models aim to learn the score function $s_\theta(x, t)$ that approximates the true gradient of the data distribution.
- It can be viewed as a probabilistic version of gradient ascent (descent)
- The original Diffusion Model is simply a discrete-time approximation of this continuous process.

Denoising Diffusion Implicit Model (DDIM)


- In a word, DDPM is “generating” at every denoising step; DDIM only generates the Gaussian noise, and the denoising steps are **fixed**.
- Assume that the variance at each denoising step is zero, and the probabilistic reverse process is determined by Expectation.
- Then, given x_t , we can compute x_0 $x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$ $x_0 = \frac{x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}}$
- So each denoising step becomes deterministic

$$x_{t-1} = \sqrt{\bar{\alpha}_{t-1}} \left(\frac{x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}} \right) + \sqrt{1 - \bar{\alpha}_{t-1}} \epsilon_\theta(x_t, t)$$

- As a result, DDIM reduces the training step from 1000 to 50.


U-Net in Diffusion Models

- U-Net consists of a U-shaped Encoder–Decoder
 - Downsample $\rightarrow$ capture context
 - Upsample $\rightarrow$ recover detail
 - Skip-connection
- Why U-Net works for Diffusion
 - Combines local (convolution) and global (up/downsampling) information.
 - Simple and stable architecture, well-suited for same-size image-to-image prediction tasks.



Stable Diffusion

- Use a pretrained VAE to map images to a low-dimensional latent, perform diffusion(add noise/denoising) there, then decode back to pixels.
- The latent space (i.e., $64 \times 64 \times 4$) is much smaller than the pixel space (i.e., $512 \times 512 \times 3$), so the training is 10 times faster.



Transformer in Diffusion Models (DiT)

- The overall diffusion process remains the same with Stable Diffusion. Only use Transformer to predict $\hat{\epsilon}_\theta(z_t, t)$
- We first apply a Convolutional VAE encoder to map the input image into a latent space z_0 .
- Then add noise to z_0 , patchify the noisy latent, and project each patch into token embeddings.
- A Transformer operates on these tokens to perform denoising in the token space.
- Then reshape back to latent space and image.

Scalable diffusion models with transformers

[PDF] thecvf.com

W Peebles, S Xie - Proceedings of the IEEE/CVF ..., 2023 - openaccess.thecvf.com

... We train latent diffusion models of images, replacing the commonly-used U-Net ... a transformer that operates on latent patches. We analyze the scalability of our Diffusion Transformers (...)

☆ 保存 引用 被引用次数: 3854 相关文章 所有 9 个版本